

NCTU RL HW2-Problem3

REINFORCE and A2C

廖家鴻 0786009
2021/5/10

Outline

- Problem3-(a): CartPole-v0 and REINFORCE
 - Network
 - Hyperparameters
 - Results
- Problem3-(b): LunarLander-v2 and A2C
 - Network
 - Hyperparameters
 - Results
 - Remarks

Problem3-(a): CartPole-v0 and REINFORCE

Prob3-(a): Network

```
45 ##### YOUR CODE HERE (5~10 lines) #####
46 ### Actor_Net ###
47 self.a_fc0 = nn.Linear(self.state_dim, self.hidden_size)
48 self.a_fc1 = nn.Linear(self.hidden_size, self.hidden_size * 2)
49 self.a_fc2 = nn.Linear(self.hidden_size * 2, self.hidden_size * 2)
50 self.a_fc3 = nn.Linear(self.hidden_size * 2, self.action_dim)
51
52 ### Baseline_Net ### (for estimating value function)
53 self.c_fc0 = nn.Linear(self.state_dim, self.hidden_size)
54 self.c_fc1 = nn.Linear(self.hidden_size, self.hidden_size)
55 self.c_fc2 = nn.Linear(self.hidden_size, self.hidden_size)
56 self.c_fc3 = nn.Linear(self.hidden_size, 1)
57 ##### END OF YOUR CODE #####
```

1. The hidden_size here is 64

```
63 def forward(self, state):
64     ##### YOUR CODE HERE (3~5 lines) #####
65     ### Actor_Net ###
66     x = F.relu(self.a_fc0(state))
67     x = F.relu(self.a_fc1(x))
68     x = F.relu(self.a_fc2(x))
69     x = F.softmax(self.a_fc3(x))
70     action_prob = x
71
72     ### Baseline_Net ###
73     y = F.relu(self.c_fc0(state))
74     y = F.relu(self.c_fc1(y))
75     y = F.relu(self.c_fc2(y))
76     y = self.c_fc3(y)
77     baseline_value = y
78     ##### END OF YOUR CODE #####
79     return action_prob, baseline_value
```

2. I implement the Actor_Net and Baseline_Net separately.

PS: I have tried another setting that is “one state input with two branches output.” And this will also work. But here I only present the separated version.

Prob3-(a): Hyperparameters

```
267 if __name__ == '__main__':  
268     # For reproducibility, fix the random seed  
269     random_seed = 20  
270     lr = 0.02  
271     env = gym.make('CartPole-v0')  
272     env.seed(random_seed)  
273     torch.manual_seed(random_seed)  
274     train(lr)  
275     test('CartPole_0.02.pth')
```

1. The learning rate is $lr=0.02$

2. The discount factor is $\gamma=0.99$

```
109 def calculate_loss(self, optimizer, gamma=0.99):  
135     g_return = self.rewards[t] + gamma*g_return
```


```
170 # Instantiate the policy model and the optimizer  
171 model = Policy()  
172 optimizer = optim.SGD(model.parameters(), lr=lr)  
173 scheduler = Scheduler.StepLR(optimizer, step_size=500, gamma=0.5)
```

3. The learning rate scheduler is set with $step_size=500$ and $\gamma=0.5$

Prob3-(a): Results

```
231 # check if we have "solved" the cart pole problem
232 if ewma_reward >= 190.:
233     torch.save(model.state_dict(), './CartPole_0.02.pth')
```

1. I save the best agent when the ewma_reward achieve 190.



2. The left figure shows the curve of ewma_reward for all training episodes.

```
Ep 1503 Length: 200 reward: 200.0000    ewma reward: 188.9586
Ep 1504 Length: 200 reward: 200.0000    ewma reward: 189.5107
Ep 1505 Length: 200 reward: 200.0000    ewma reward: 190.0352
Solved! Running reward is now 190.03516076169097 and the last episode runs to 200 time steps!
Episode 1   Reward: 200.0
Episode 2   Reward: 200.0
Episode 3   Reward: 200.0
Episode 4   Reward: 200.0
Episode 5   Reward: 200.0
Episode 6   Reward: 200.0
Episode 7   Reward: 200.0
Episode 8   Reward: 200.0
Episode 9   Reward: 200.0
Episode 10  Reward: 200.0
```

3. At this setting, the total training episodes are 1505. And the last episode runs to 200 steps.

4. I follow the sample code to test the trained agent for 10 episodes and all get the maximum reward.

Problem3-(b): LunarLander-v2 and A2C

Prob3-(b): Network

```
43 self.hidden_size = 64
44 ##### YOUR CODE HERE (5~10 lines) #####
45 self.fc0 = nn.Linear(self.state_dim, self.hidden_size)
46 ### Actor_Net ###
47 self.a_fc0 = nn.Linear(self.state_dim, self.hidden_size)
48 self.a_fc1 = nn.Linear(self.hidden_size, self.hidden_size * 2)
49 self.a_fc2 = nn.Linear(self.hidden_size * 2, self.hidden_size * 2)
50 self.a_fc3 = nn.Linear(self.hidden_size * 2, self.action_dim)
51
52 ### Baseline_Net ### (for estimating value function)
53 self.c_fc0 = nn.Linear(self.state_dim, self.hidden_size)
54 self.c_fc1 = nn.Linear(self.hidden_size, self.hidden_size)
55 self.c_fc2 = nn.Linear(self.hidden_size, self.hidden_size)
56 self.c_fc3 = nn.Linear(self.hidden_size, 1)
57 ##### END OF YOUR CODE #####
```

1. The hidden_size here is 64.

```
63 def forward(self, state):
64     ##### YOUR CODE HERE (3~5 lines) #####
65     # s = F.relu(self.fc0(state))
66     ### Actor_branch ###
67     x = F.relu(self.a_fc0(state))
68     x = F.relu(self.a_fc1(x))
69     x = F.relu(self.a_fc2(x))
70     x = F.sigmoid(self.a_fc3(x))
71     action_prob = x
72
73     ### Baseline_branch ### (for estimating value function)
74     y = F.relu(self.c_fc0(state))
75     y = F.relu(self.c_fc1(y))
76     y = F.relu(self.c_fc2(y))
77     y = self.c_fc3(y)
78     baseline_value = y
79     ##### END OF YOUR CODE #####
80     return action_prob, baseline_value
```

2. Again, I implement the Actor_Net and Baseline_Net separately.

PS: I also tried another setting that is “one state input with two branches output” , but it is hard work.

Prob3-(a): Hyperparameters

```
261 if __name__ == '__main__':  
262     # For reproducibility, fix the random seed  
263     random_seed = 20  
264     lr = 0.01  
265     env = gym.make('LunarLander-v2')  
266     env.seed(random_seed)  
267     torch.manual_seed(random_seed)  
268     train(lr)  
269     test('LunarLander_0.01.pth')
```

1. The learning rate is $lr=0.01$

```
103 def calculate_loss(self, gamma=0.95):  
132     baseline_loss = nn.MSELoss()  
133     value_loss = baseline_loss(value, reward+(gamma*next_value) )  
134  
135     advantage = (reward + gamma*next_value) - value  
136     policy_loss = -log_act_prob * advantage.detach()
```


2. The discount factor is $gamma=0.95$

```
160 # Instantiate the policy model and the optimizer  
161 model = Policy()  
162 optimizer = optim.SGD(model.parameters(), lr=lr)  
163 # optimizer = optim.Adam(model.parameters(), lr=lr)  
164 scheduler = Scheduler.StepLR(optimizer, step_size=100, gamma=0.8)
```

3. The learning rate scheduler is set with $step_size=100$ and $gamma=0.8$

Prob3-(b): Results

```
222 # check if we have "solved" the cart pole problem
223 if ewma_reward >= 172.:
224     torch.save(model.state_dict(), './LunarLander_0.01.pth')
225     print("Solved! Running reward is now {} and "
226           "the last episode runs to {} time steps!".format(ewma_reward, t))
227     break
```



1. I save the best agent when the ewma_reward achieve 172.

2. The left figure shows the curve of ewma_reward for all training episodes.

```
Ep 230 Length: 516 R: 200.5952 ewma reward: 162.8868
Ep 231 Length: 267 R: 247.0709 ewma reward: 167.0960
Ep 232 Length: 395 R: 233.3892 ewma reward: 170.4107
Ep 233 Length: 870 R: 212.3251 ewma reward: 172.5064
Solved! Running reward is now 172.506412644235 and the last episode runs to 870 time steps!
Episode 1 Reward: 247.8139565405427
Episode 2 Reward: 277.3608585658985
Episode 3 Reward: -1.5739138103989347
Episode 4 Reward: 202.5240771279636
Episode 5 Reward: 244.86838236388374
Episode 6 Reward: 235.36878862993836
Episode 7 Reward: 141.4510078744156
Episode 8 Reward: 239.24533874901292
Episode 9 Reward: 74.15252161222428
Episode 10 Reward: -191.0576380074957
```

3. At this setting, the total training episodes are 233.

4. I follow the sample code to test the trained agent for 10 episodes, only episode 3 and 10 are failedg.

Prob3-(b): Remarks

```
##### YOUR CODE HERE (10-15 lines) #####
for t in itertools.count(start=1):
    action = model.select_action(state)
    next_state, reward, done, _ = env.step(action)

    ep_reward += reward
    state = next_state
    model.rewards.append(reward)

    loss, policy_loss, value_loss = model.calculate_loss(reward, next_state)
    # optimizer.zero_grad()
    # loss.backward()
    # nn.utils.clip_grad_norm_(model.parameters(), 3)
    # optimizer.step()

    ### Update Actor_Net ###
    optimizer.zero_grad()
    policy_loss.backward(retain_graph=True)
    nn.utils.clip_grad_norm_(model.parameters(), 3)
    optimizer.step()

    ### Update Behavior_Net ###
    optimizer.zero_grad()
    value_loss.backward()
    nn.utils.clip_grad_norm_(model.parameters(), 3)
    optimizer.step()

    if done:
        break

model.clear_memory()
##### END OF YOUR CODE #####
```

The key point that I finally success is that the network update frequency.

Originally, I update the network per episode as 3-(a), but it always fail no matter I tune the hyperparameters.

After I change the way to update the network by every step, the agent finally works!!!

The End